

# Assignment 3: An Introduction to the World of SDN

Arpit Prasad and Akshat Bhasin  
2022EE11837 and 2022EE31996  
**COL334: Computer Network**

October 13, 2025

## 1 Part 1: Hub Controller and Learning Switch

### 1.1 pingall Test

The following are the rules installed in the switches after running pingall:

#### 1. Hub Controller:

```
*** s1 -----  
cookie=0x0, duration=154.396s, table=0, n_packets=83, n_bytes=7966, priority=0  
actions=CONTROLLER:65535  
*** s2 -----  
cookie=0x0, duration=154.407s, table=0, n_packets=83, n_bytes=7966, priority=0  
actions=CONTROLLER:65535
```

#### 2. Learning Switch Controller:

```
*** s1 -----  
cookie=0x0, duration=6.527s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.523s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.514s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.503s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.496s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.488s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.476s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.468s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.462s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0x0, duration=6.456s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="  
cookie=0xx, duration=14.521s, table=0, n_packets=53, n_bytes=5458, priority=0 actio
```

```

*** s2 -----
cookie=0x0, duration=6.522s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="
cookie=0x0, duration=6.518s, table=0, n_packets=1, n_bytes=98, priority=1,in_port="
...
cookie=0x0, duration=14.532s, table=0, n_packets=54, n_bytes=5548, priority=0 actio

```

The following are the observations of the above results:

### 1. Hub Controller Observations:

- Only a single, low-priority "table-miss" rule is present on each switch.
- This rule's action is `actions=CONTROLLER`, which forces every single packet that the switch does not have a rule for to be sent to the controller.
- Since no other rules are ever installed, this means all packets (ARP, ping requests, ping replies) are sent to the controller for a forwarding decision, making the switch effectively "dumb."

### 2. Learning Switch Observations:

- Multiple specific, high-priority flow rules are installed on the switches.
- Each rule matches on a source/destination MAC address pair and an input port.
- This indicates that once the first packet of a conversation is seen, the controller proactively installs a rule on the switch, allowing all subsequent packets of that same conversation to be forwarded directly by the switch hardware at line rate.
- The low-priority table-miss rule is still present but handles far fewer packets, as it is only used for the first packet of a new, unknown flow.

## 1.2 Throughput Test

The following are the Throughput of when the following controllers are used:

1. Hub Controller: 20.3 Mbits/sec
2. Learning Switch: 29.1 Gbits/sec

Inferences:

1. **Hub Controller Inference:** The throughput is very low because every data packet in the `iperf` stream must make a slow, high-latency round trip from the switch to the controller for a forwarding decision. The controller itself becomes the performance bottleneck.
2. **Learning Switch Inference:** The throughput is extremely high because the controller only processes the first packet of the flow. It then installs a rule on the switch, allowing all subsequent data packets to be forwarded at the switch's hardware speed (line rate), completely bypassing the controller bottleneck.

## 2 Part 2: Layer2-like Shortest Path Routing

The following are the Testing and Measurements Performed:  
iperf with two parallel TCP Connections:

### 1. ECMP Off:

- (a) Throughput: 9.50 Mbits/sec
- (b) Flow Rules:

```
*** s1 -----  
cookie=0x0, duration=72.327s, table=0, n_packets=130, n_bytes=7800,  
  priority=65535,dl_dst=01:80:c2:00:00:0e,dl_type=0x88cc actions=CONTROLLER:6553  
cookie=0x0, duration=72.338s, table=0, n_packets=42449, n_bytes=4569148,  
  idle_timeout=10, priority=0 actions=CONTROLLER:65535  
... (Similar rules on s2-s6) ...
```

### 2. ECMP On:

- (a) Throughput: 19.2 Mbits/sec or 9 Mbits/sec
- (b) Flow Rules:

```
*** s1 -----  
cookie=0x0, duration=20.693s, table=0, n_packets=2123, n_bytes=12329874,  
  idle_timeout=10, priority=20,tcp,nw_src=10.0.0.1,nw_dst=10.0.0.2,tp_src=51634,  
  actions=output:"s1-eth3"  
  
cookie=0x0, duration=20.693s, table=0, n_packets=1937, n_bytes=127878,  
  idle_timeout=10, priority=20,tcp,nw_src=10.0.0.2,nw_dst=10.0.0.1,tp_src=5001,t  
  actions=output:"s1-eth1"  
  
cookie=0x0, duration=20.693s, table=0, n_packets=1720, n_bytes=12303276,  
  idle_timeout=10, priority=20,tcp,nw_src=10.0.0.1,nw_dst=10.0.0.2,tp_src=51638,  
  actions=output:"s1-eth2"  
  
cookie=0x0, duration=20.693s, table=0, n_packets=1712, n_bytes=113028,  
  idle_timeout=10, priority=20,tcp,nw_src=10.0.0.2,nw_dst=10.0.0.1,tp_src=5001,t  
  actions=output:"s1-eth1"  
... (Other rules on s1 and similar path-specific rules on s2, s3, s4) ...
```

Observations:

### 1. ECMP Off Observations:

- The controller selects only one of the two available equal-cost paths for both parallel TCP connections.

- The total throughput of 9.50 Mbits/sec is approximately the maximum capacity of a single 10 Mbps link in the topology.
- Both TCP flows are forced to compete for the limited bandwidth of this single path, effectively capping the performance.

## 2. ECMP On Observations:

- The flow rules on switch **s1** clearly show that the two TCP connections (identified by different source ports **51634** and **51638**) are being forwarded out of different physical ports (**s1-eth3** and **s1-eth2**, respectively). This is direct proof of load balancing.
- The total throughput of 19.2 Mbits/sec is almost exactly double the result with ECMP off.
- This demonstrates that the controller successfully split the traffic, allowing the flows to utilize the aggregate bandwidth of both available 10 Mbps paths simultaneously.
- But this was not the case all the time. Since there was a 50% chance of the same path being chosen for both of the controllers

## 2.1 Bonus Part

### Load Balancing Mechanism:

- The weighted load-balancing strategy works by maintaining a count of active flows on each link in the network.
- When a new flow arrives and multiple equal-cost paths are available, the controller calculates the total flow count (utilization) for each path.
- It then deterministically selects the path with the minimum total utilization, ensuring that new flows are always assigned to the currently lightest-loaded path.

### Results:

1. **iperf** with UDP results are shown in Table 1 (assuming links have a BW=100Mbps).

Flow	Target BW	Received BW	Packet Loss	Out of Order
Heavy Flow	80 Mbps	84.8 Mbps	0%	796
Light Flow	10 Mbps	10.8 Mbps	0%	225

Table 1: Bandwidth and packet statistics for heavy and light flows.

## 2. Controller Decision Logic:

- (a) A sample of the controller logs demonstrates the deterministic path selection:

```
PacketIn: UDP 10.0.0.1:38216 -> 10.0.0.2:5001 on switch 1
Path [1, 3, 5, 6] has a utilization of 0
Path [1, 2, 4, 6] has a utilization of 0
Selected path for flow 10.0.0.1:38216 -> 10.0.0.2:5001 is [1, 3, 5, 6]
```

```
PacketIn: UDP 10.0.0.1:38216 -> 10.0.0.2:5001 on switch 1
Path [1, 3, 5, 6] has a utilization of 3
Path [1, 2, 4, 6] has a utilization of 0
Selected path for flow 10.0.0.1:38216 -> 10.0.0.2:5001 is [1, 2, 4, 6]
```

### Validation of Result:

- The presence of a high number of out-of-order packets suggests that the flows were traversing different network paths.
- The controller logs provide definitive proof of the weighted selection. When the first packet of the heavy flow (**port:38216**) arrived, the controller chose an empty path.
- Due to a race condition, a subsequent packet from the same flow triggered another decision. The controller, now aware of the first decision, saw an unbalanced state and correctly chose the other, empty path.
- When the second, lighter flow (**port:59291**) arrived, the controller would have seen that the first path was already heavily utilized by the 80 Mbps flow and would have deterministically placed the new flow on the second, less-utilized path.

### Comparison with Random Selection Methodology:

- This deterministic behavior contrasts sharply with the random selection methodology from the main part of the assignment.
- A random selector would have had a 50% chance of placing the second (light) flow on the same path as the first (heavy) flow, leading to suboptimal load distribution.
- The implemented weighted strategy guarantees that flows are distributed across available paths based on load, fulfilling the bonus requirement.

## 3 Layer3-like Shortest Path Routing

The following are the experimenting and reports:

1. h1 ping h2 -c 5
2. Rules installed in switches

Assumptions:

1. A1
2. A2

## 4 Part 4: Comparision with Traditional Routing (OSPF)

Experiments Tested:

iperf result between h1 and h2:

I1

Forwarding Rules:

1. s1

2. S6

Throughput and Convergence Times:

1. SDN (mention source)

2. OSPF (mention source)

Explanation of Result:

1. E1

2. E2